

Exercise 2015-09-21

Hubert Rehrauer

September 21, 2015

1 Introduction

A study examines a liver disease. In affected patients the liver does not work properly. Part of the liver shows acute signs other parts are affected but do not show severe symptoms. The study searches for the genes that have changed expression due to the disease. The goal is to understand at a molecular level the causes and consequences of the disease.

The study has used Affymetrix Whole Genome Microarrays of the type HG-U133-Plus2 to measure the gene expression. In total 5 sick patients were examined and from each patient a sample from acutely and moderately tissue was measured. As a reference normal tissue from 6 healthy patients were examined.

In this exercise we run some exploratory analysis of the data set. We want to check if there are outliers and/or systematic biases.

2 Loading the data

The phenotype information is in the file `SampleAnnotation.txt` that we read into a data frame,

```
> anno = read.table("SampleAnnotation.txt", as.is=TRUE, sep="\t", quote="",  
+                  row.names=1, header=TRUE)
```

that holds the sample name, tissue type, patient ID, and associated file:

```
> anno
```

	TissueType	PatientId	File
norm-02	norm	P02	norm-02.CEL
norm-05	norm	P05	norm-05.CEL
norm-07	norm	P07	norm-07.CEL
norm-09	norm	P09	norm-09.CEL
norm-10	norm	P10	norm-10.CEL
norm-11	norm	P11	norm-11.CEL
sick-04	sick	P04	sick-04.CEL
sick-12	sick	P12	sick-12.CEL
sick-13	sick	P13	sick-13.CEL
sick-14	sick	P14	sick-14.CEL
sick-15	sick	P15	sick-15.CEL

For labeling and coloring the subsequent plots we define the variables:

```
> samples = rownames(anno)
> colors = rainbow(nrow(anno))
```

and generate boolean indices with which we can access the normal, sick and acute samples only:


```
> isNorm = anno$TissueType == "norm"
> isSick = anno$TissueType == "sick"
> isAcute = anno$TissueType == "acute"
```

Now we load the expression data

```
> x = read.table("expressiondata.txt", as.is=TRUE, sep="\t", quote="", row.names=1, header=1)
> x = as.matrix(x)
```

We visually compare the expression signals from sample 1 and 2 by plotting them with the standard plot command:

```
> plot(x[, "norm-02"], x[, "norm-05"], log="xy", pch=20)
> abline(0, 1, col="blue")
> abline(log10(2), 1, col="blue", lty=2);
> abline(-log10(2), 1, col="blue", lty=2);
```



The solid blue line gives the first diagonal and the dashed lines give the boundaries for 2-fold up- or down-regulation.

3 Checking the distribution of the intensities

A basic assumption for the subsequent analysis is, that the intensity distributions of the different arrays are similar. Do create boxplots that summarize the value distribution for each sample, and use the function `plotDensities` in the package `limma`.

4 Checking the consistency of the replicates

Do check the consistency of the replicates by computing sample correlation.

In order to check whether the replicates correlate well, we compute the correlation matrix on the logarithmic values and print it


```
> corrMatrix = cor(x)
> signif(corrMatrix, digits=3)
```

	norm-02	norm-05	norm-07	norm-09	norm-10	norm-11	sick-04	sick-12	sick-13
norm-02	1.000	0.981	0.962	0.967	0.958	0.982	0.910	0.885	0.882
norm-05	0.981	1.000	0.963	0.978	0.975	0.980	0.938	0.927	0.916
norm-07	0.962	0.963	1.000	0.984	0.969	0.968	0.940	0.925	0.925
norm-09	0.967	0.978	0.984	1.000	0.985	0.970	0.954	0.945	0.938
norm-10	0.958	0.975	0.969	0.985	1.000	0.974	0.955	0.950	0.939
norm-11	0.982	0.980	0.968	0.970	0.974	1.000	0.926	0.905	0.901

sick-04	0.910	0.938	0.940	0.954	0.955	0.926	1.000	0.983	0.976
sick-12	0.885	0.927	0.925	0.945	0.950	0.905	0.983	1.000	0.974
sick-13	0.882	0.916	0.925	0.938	0.939	0.901	0.976	0.974	1.000
sick-14	0.957	0.980	0.973	0.990	0.988	0.968	0.962	0.960	0.946
sick-15	0.862	0.886	0.930	0.932	0.940	0.897	0.939	0.943	0.938
	sick-14	sick-15							
norm-02	0.957	0.862							
norm-05	0.980	0.886							
norm-07	0.973	0.930							
norm-09	0.990	0.932							
norm-10	0.988	0.940							
norm-11	0.968	0.897							
sick-04	0.962	0.939							
sick-12	0.960	0.943							
sick-13	0.946	0.938							
sick-14	1.000	0.936							
sick-15	0.936	1.000							

We visualize the matrix as an image:

```
> par(mar=c(8,8,2,2))
> grayScale <- gray((1:256)/256)
> image(corrMatrix, col=grayScale, axes=FALSE)
> axis(1, at=seq(from=0, to=1, length.out=length(samples)), labels=samples, las=2)
> axis(2, at=seq(from=0, to=1, length.out=length(samples)), labels=samples, las=2)
```




From the correlation we can see that

- normals show high correlation among each other
- the normal are very different from both sick and acute
- the sick and acute are rather similar
- sick-14 rather looks like a normal sample
- sick-15 has overall low correlation but rather looks like an "acute"
- acute-04-a which is a technical replicate of acute-04 is more similar to sick-04 than to acute-04.

5 Sample Clustering


The sample clustering shows the similarities of the expression patterns of the samples in a tree. In order to compute the similarities of the samples one can use all genes or only a subset of the genes. When using all genes where is the risk that the absent genes drive the clustering of the samples. This is because in many studies the absent genes make up the majority of the measured genes. But those genes have a low intensity that is strongly influenced by the background signal measured on each chip and not by real gene expression

```
> x.sd = apply(x, 1, sd, na.rm=TRUE)
> ord = order(x.sd, decreasing=TRUE)
> highVarGenes = ord[1:500]
> par(mfrow=c(1,2));
> d = as.dist(1-cor(x));
> c=hclust(d, method="ward.D2");
> plot(c, hang=-0.1, main="all genes", xlab="")
> d = as.dist(1-cor(x[highVarGenes, ]));
> c=hclust(d, method="ward.D2");
> plot(c, hang=-0.1, main="high variance genes", xlab="")
```



If we run the clustering without sample sick-04, the acute-04 does no longer cluster in the branch with the other sick samples

```
> sub = x[, samples != "sick-04"]
> d = as.dist(1-cor(sub));
> c=hclust(d, method="ward.D2");
> plot(c, hang=-0.1, main="all genes", xlab="")
```



6 Apply quantile normalization

Use the function `limma::normalizeQuantiles` to normalize the data and plot again the histograms.

7 Sample Representation in Principal Component Space

Use the functions `cmdscale` and `prcomp` to create a plot that represents the sample distances in a reduced space.